

# Memcached: The Key-Value Architecture that Revolutionized Distributed Storage in Big Data

João Santos<sup>1</sup>

<sup>1</sup>Department of Computer Engineering, ISEC, Rua Pedro Nunes - Quinta da Nora, Coimbra, 3030-199, Portugal, <https://www.isec.pt/>

\*João Santos.

## Abstract

Memcached, a high-performance distributed in-memory key-value store, has emerged as a critical tool for optimizing web applications and Big Data workloads by reducing latency and offloading backend database pressure. This paper explores Memcached's architecture, implementation, and practical applications in modern distributed systems.

**Keywords:** Memcached, Big Data, Key-Value Store, Performance.

## Nomenclature

### 1. Introduction

This paper will begin by contextualizing Memcached role in the NoSQL ecosystem, detailing its slab-based memory management, stateless design, and reliance on volatile RAM for ultra-fast operations. A step-by-step installation guide is provided, along with an analysis of its write-through caching mechanism, where data is first stored in Memcached to accelerate subsequent requests. Challenges such as data volatility (no native redundancy) and item size limitations (1MB default) are discussed, along with workarounds like integrating Redis for persistent storage. Real-world applications are highlighted, including Facebook's use of Memcached for session caching, Twitter's optimization of popular post retrieval, and YouTube's latency reduction in video recommendation systems. A critical comparison with Redis, its closest competitor, concludes the analysis. While Redis offers advanced data structures (e.g., lists, sets) and persistence, Memcached excels in simplicity and raw speed for pure caching scenarios. Performance benchmarks highlight Memcached's strengths in low-latency scenarios, while scalability tests reveal trade-offs in horizontal expansion.

## 2. The Origins of Memcached 30

Memcached was originally developed by Brad Fitzpatrick in 2003 for his social networking website, LiveJournal. LiveJournal had to deal with serious scalability problems at the time due to increasing user populations and increasing complexity of database queries. Fitzpatrick realized that many database queries were repetitive and could be optimized by caching frequently accessed data in memory. This led to the creation of Memcached as a solution to reduce database load and improve application performance.

The name "Memcached" is a portmanteau of "memory" and "cache daemon", reflecting its core functionality as a memory-based caching system. Since its release, Memcached has become a crucial resource for developers and organizations interested in maximizing the performance of their web applications.

## 3. Evolution and Adoption of Memcached 41

Since its creation in 2003, Memcached has been under significant improvements and has been widely adopted by major tech companies. Some of this notable milestones were included in:

**2004:** Memcached was open-sourced, allowing developers worldwide to contribute to its development and use it in their projects.

**2008:** Facebook, one of the largest users of Memcached, shared information into how they scaled Memcached to handle billions of requests of users per day. Nishtala et al., 2008

**2010s:** Memcached became a standard component in the tech stacks of companies like Twitter, YouTube, and Wikipedia, further solidifying its reputation as a reliable caching solution for the company.

Over the years, Memcached has also given rise to other caching systems and technologies, such as Redis, which extends the concept of in-memory caching with additional features like persistence and advanced data structures. Memached, 2023

## 4. Memcached role in the NoSQL Ecosystem 56

NoSQL (Not Only SQL) databases emerged as a response to the limitations of traditional relational databases (RDBMS) in handling modern web-scale applications. Unlike RDBMS, which enforce rigid schemas and transactional consistency (ACID properties), NoSQL systems prioritize scalability and high availability, particularly in distributed environments. NoSQL databases are categorized into four primary types:

**Key-Value:** Simple data models where each item is accessed via a unique key (e.g., Memcached, Redis).

**Document:** Store semi-structured data like JSON or XML (e.g., MongoDB).

**Column-Family:** Optimized for querying large datasets (e.g., Cassandra). 65  
**Graph:** Designed for interconnected data (e.g., Neo4j). 66  
Memcached, as a key-value store, emphasizes the NoSQL philosophy by sacrificing 67  
complex querying and persistence for raw speed and horizontal scalability. 68

## 4.1 Memcached Role in the NoSQL environment 69

Memcached operates fundamentally differently from traditional relational databases or 70  
conventional NoSQL databases (such as MongoDB or Redis). Instead, it functions as a 71  
high-performance, distributed memory caching system designed to reduce database load 72  
and increase application performance, which is the main goal. By temporarily storing 73  
frequently accessed data such as query results, session states, or computational objects 74  
in RAM, Memcached reduces latency and reduces the repetitive read operations from 75  
backend databases. This makes it a critical complementary layer for both SQL databases 76  
and NoSQL databases. 77

Its architectural philosophy aligns three critical aspects of NoSQL principles: 78

### 4.1.1 Scalability-First Design: 79

Like NoSQL systems, Memcached prioritizes horizontal scaling over vertical scaling. It 80  
distributes data across a cluster of nodes using consistent hashing, allowing seamless 81  
expansion by adding servers without downtime or complex sharding logic. This contrasts 82  
with traditional databases, which often require disruptive scaling procedures. 83

### 4.1.2 Schema-Less Data Model: 84

Memcached adopts a NoSQL-like key-value store structure, where data is stored as un- 85  
structured blobs identified by unique keys. Unlike relational databases, it does not enforce 86  
rigid schemas or relationships, allowing rapid writes and flexibility for caching dynamic 87  
or heterogeneous data. 88

### 4.1.3 Transitory Performance-Oriented Operations: 89

Mirroring NoSQL's focus on speed and simplicity, Memcached sacrifices durability fea- 90  
tures (e.g., disk persistence) to maximize throughput. It operates entirely in memory, 91  
with optional LRU (Least Recently Used) eviction policies, prioritizing low-latency access 92  
over guarantees of data permanence, a trade-off common in NoSQL systems like Apache 93  
Cassandra when configured for high-speed writes. 94

However, Memcached diverges from NoSQL databases in key areas: It lacks native 95  
replication, query languages, or persistence mechanisms, instead relying on external data- 96  
bases as the main source of data. This reinforces its specialized role as a transient caching 97  
layer rather than a standalone data storage solution. 98

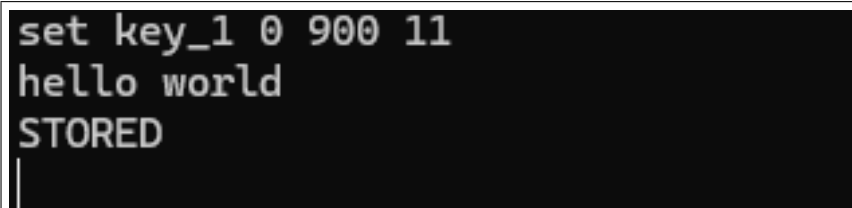
## 4.2 Key-Value Stores: The Foundation of Memcached Explained 99

Key-value stores are the simplest and most performant category of NoSQL databases. 100

Memcached exemplifies this model through: 101

### 4.2.1 Core Operations 102

**SET(key, value):** Store a value with an expiration time. 103



```
set key_1 0 900 11
hello world
STORED
|
```

Figure 1: Command line insert operation test

**key\_1:** This is the name of the key where the value will be stored. In the example, 104  
the key is key\_1. This key can be used to retrieve the stored value later. 105

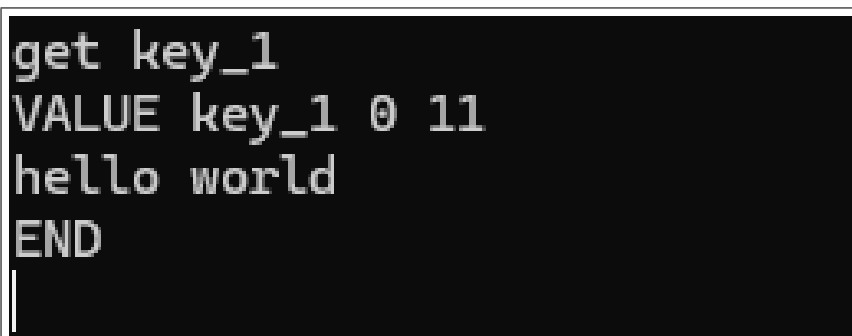
**0:** This is the flags field. It is an integer that can be used to store additional metadata 106  
about the value. In this case, the value is 0, which means no specific flag has been set. 107

**900:** This is the expiration time in seconds. The value 900 means that the item will 108  
expire and be removed from the cache after 900 seconds (15 minutes). If the value is 0, 109  
the item never expires. 110

**11:** This is the size in bytes of the value to be stored. In this example, the value hello 111  
world is 11 bytes long (including the space). 112

**hello world:** This is the value that will be stored in the key key\_1. In this example, 113  
the value is the string hello world. 114

**GET(key):** Retrieve a value by its key. 115



```
get key_1
VALUE key_1 0 11
hello world
END
|
```

Figure 2: Command line get operation test

**key\_1:** This is the name of the key where the value was stored. In this example, the 118  
key is key\_1. You can use this key to retrieve the stored value later. 119

120

**DELETE(key):** Remove a key-value pair. 121  
122



Figure 3: Command line delete operation test

**key\_1:** This is the name of the key where the value was stored. In this example, the 123  
key is key\_1. You can use this key to retrieve the stored value later. 124

### 4.3 Use Cases in NoSQL Workflows 125

#### 4.3.1 Session Caching: Storing User Sessions in Memory to Reduce Database 126 Queries 127

In contemporary web applications, efficient handling of user sessions is crucial to giving 128  
a responsive and smooth experience. Session data like login status, user preferences, or 129  
items in a shopping cart would typically be held in databases. However, retrieving from 130  
the database for each and every session-related operation can turn the process slow and 131  
add additional load on the system. 132

That's where Memcached comes in. It's used as a super-fast, in-memory session cache. 133  
When a user logs in or interacts with the application, Memcached caches the session data 134  
with a session ID as the key. From then on, future requests for this session data are pulled 135  
from memory and not visited by the database at all. 136

Imagine an e-commerce site in the middle of a flash sale. With millions of shoppers 137  
browsing and shopping, Memcached can store session data for all of them, enabling rapid 138  
page loads and a seamless shopping experience, even when traffic spikes. 139

#### 4.3.2 Query Result Caching: Accelerating Repeated Database Queries 140

Many applications, especially those with high read-to-write ratios, frequently execute the 141  
same database queries. For instance, an online store might repeatedly fetch product 142  
listings, category filters, or search results. Executing these queries repeatedly can strain 143  
the database and slow down the application. 144

Memcached solves this problem by caching the results of frequently executed queries. 145  
When a query is first run, its result is stored in Memcached with a unique key (e.g., a 146  
hash of the query parameters). Subsequent requests with the same parameters retrieve 147  
the cached result directly from memory, eliminating the need to re-execute the query. 148

For example, a news website might cache the results of popular queries like "top 149  
headlines" or "trending articles," ensuring fast delivery of content to millions of readers. 150

### 4.3.3 Temporary Data Storage 151

The majority of applications store and use transient or ephemeral data that doesn't need 152  
to be stored forever in a database. API keys, rate limiting counts, or temporary state 153  
data for background jobs are some examples. Storing it permanently in a conventional 154  
database might be wasteful and inefficient, considering it will introduce overhead and 155  
consume valuable storage space. 156

Memcached is ideally suited to cache such temporary data since it is in-memory and 157  
lightweight. It provides fast read/write access and automatically evicts data when no 158  
longer needed (e.g., on TTL expiration or memory pressure). This makes it ideally suited 159  
for use cases like: 160

**API token caching:** Temporarily caching OAuth tokens or session tokens to au- 161  
thenticate API calls without having to query an authentication service repeatedly. 162

**Rate limiting:** Tracking the number of requests by a user or IP address within a 163  
time window to cap usage. 164

**Temporary state storage:** Keeping intermediate results of background tasks (e.g., 165  
batch processing or data aggregation) until the task is complete. 166

For example, a social networking site can use Memcached to cache transient counters 167  
for the number of shares or likes on a post in real-time, allowing for instant update without 168  
putting too much burden on the database. 169

## 5. Technical Details of Memcached 170

### 5.1 Architecture Overview 171

Memcached is designed as a distributed in-memory caching system, optimized for high- 172  
speed data access. In the bellow section will be explained the following key components: 173  
174

**Distributed Nature:** Memcached operates as a cluster of servers, each holding a 175  
portion of the cached data. 176

Data is distributed across nodes using a consistent hashing algorithm, ensuring min- 177  
imal rehashing when nodes are added or removed. 178

**Stateless Design:** Memcached servers are stateless, meaning they do not communic- 179  
ate with each other directly. 180

Clients are responsible for determining which server stores a specific key using the 181  
hashing algorithm. 182

**Client-Server Model:** Clients (ex: web applications) communicate with Memcached 183  
servers via a simple TCP/IP or UDP protocol. 184

There are various libraries that are available in multiple programming languages (Py- 185  
thon, PHP, Java, etc.) to facilitate integration of Memcached in the business. 186

## 5.2 Memory Management in Memcached

187

Memcached's memory management is one of its most distinctive features, designed to maximize efficiency and minimize fragmentation:

188

189

### 5.2.1 Cases of Fragmentation

190

Memory fragmentation occurs when free memory spaces are divided into small, non-contiguous segments of memory so that it is difficult to assign large portions of memory even though the free space is abundant. It can happen in two forms: external fragmentation (free memory divided into small chunks) and internal fragmentation (free memory within occupied blocks). It can slow down the performance of the system and lead to memory wastage. In the picture bellow we can see one example of this problem that Memcached can fix.

191

192

193

194

195

196

197




Figure 4: Memory Fragmentation Example

### 5.2.2 Page Allocation:

198

To fix the problem explained above Memcached came with an implementation where memory is divided into pages, which are divided providing chunks of memory categorized by size according to the slab class which they belong too.

199

200

201




Figure 5: Visual Example of Memory Allocation in Memcached

### 5.2.3 Slab Classes:

In Memcached, memory is managed using a system of pages, which are pre-allocated chunks of memory divided into smaller, fixed-size blocks. These pages are grouped into slab classes, with each class responsible for handling items within a specific size range. For example, Slab Class 1 might manage items up to 72 bytes, while Slab Class 43 could handle items as large as 1MB.

The division of slab classes is designed to optimize memory usage and minimize waste. Each slab class is tailored to a specific size range, ensuring that items are stored in the smallest possible block that can accommodate them. This prevents scenarios where a large block is used to store a small item, which would lead to internal fragmentation (unused memory within a block).

#### How Slab Classes Are Calculated:

Memcached uses a growth factor to determine the size ranges for each slab class. The growth factor is a configurable parameter (default is 1.25) that defines how quickly the block sizes increase from one slab class to the next. The size of each slab class is calculated using the formula:

$$\text{Size of Slab Class } n = \text{Base Size} \times (\text{Growth Factor})^{n-1}$$

For example, if the base size is 64 bytes and the growth factor is 1.25, the size of Slab Class 1 would be 64 bytes, Slab Class 2 would be 80 bytes ( $64 \times 1.25$ ), Slab Class 3 would be 100 bytes ( $80 \times 1.25$ ), and so on.



#### 5.2.4 Memory ejection Policy:

221

Items are evicted if they have not expired (an expiration time of 0 or some time in the future), the slab class is completely out of free chunks, and there are no free pages to assign to a slab class.

222  
223  
224

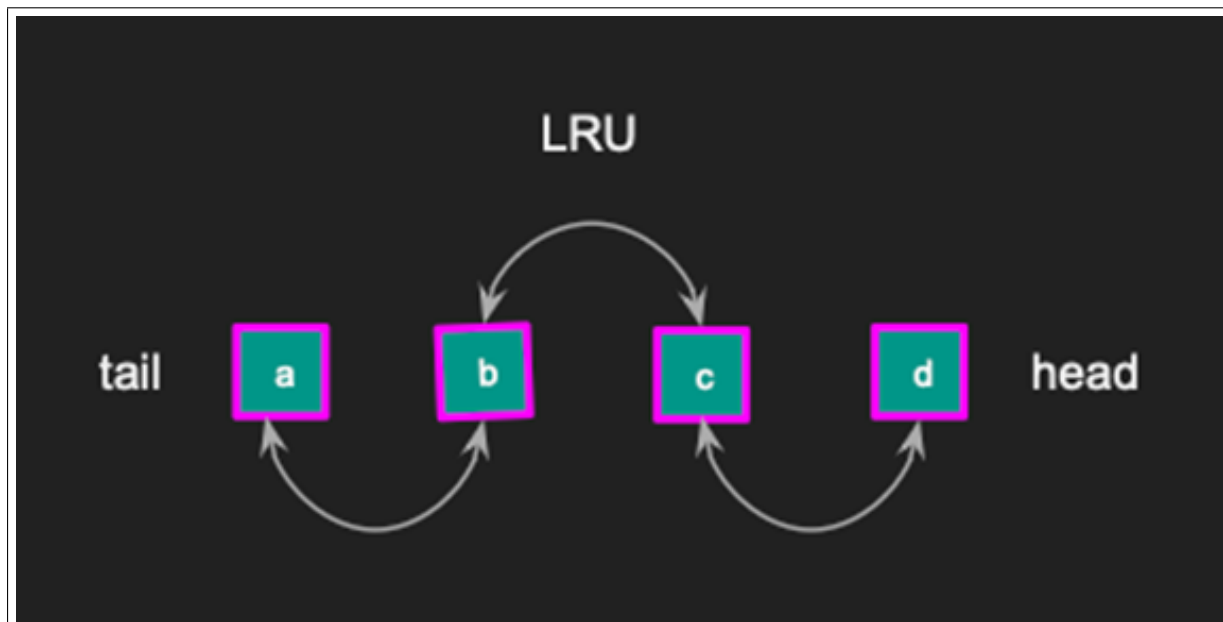


Figure 6: LRU Usage Example

Memory is also reclaimed when it's time to store a new item. If there are no free chunks, and no free pages in the appropriate slab class, memcached will look at the end of the LRU for an item to “reclaim”. It will search the last few items in the tail for one which has already been expired, and is thus free for reuse. If it cannot find an expired item however, it will “evict” one which has not yet expired. This is then noted in several statistical counters. Memcached, 2023

225  
226  
227  
228  
229  
230

### 5.3 Operating System Compatibility

231

Memcached is highly versatile and runs on a variety of operating systems, making it suitable for diverse environments:

232  
233

#### **Linux:**

234

Memcached is natively supported on most Linux distributions (Ubuntu, CentOS, Debian, etc.).

235  
236

Installation is straightforward via package managers like apt or yum.

237

#### **Windows:**

238

While not natively supported, Memcached can be run on Windows using the Windows Subsystem for Linux (WSL) or pre-compiled binaries.

239  
240

#### **macOS:**

241

Memcached can be installed via package managers like Homebrew or compiled from source.

242  
243

<b>5.4 Hardware Requirements</b>	244
Memcached is for default lightweight but has some specific hardware requirements for optimal performance:	245 246
<b>5.4.1 RAM:</b>	247
RAM is the most critical resource for Memcached, as it stores all data in memory. The amount of RAM required depends on the size of the cache you want to maintain.	248 249
<b>Minimum Recommended:</b>	250
<b>512MB:</b> Suitable for development environments or small testing setups.	251
<b>1GB:</b> For production environments with light to moderate loads.	252
<b>4GB or more:</b> For production environments with high loads or large data volumes.	253 254
<b>How to Calculate:</b> Estimate the total size of the data you want to store in the cache.	255 256
Add a safety margin (20-30%) to avoid running out of memory.	257
<b>Example:</b> If the system needs to store 2GB of data in the cache, allocate at least 2.5GB of RAM.	258 259
<b>Configuration:</b> The amount of memory allocated to Memcached is defined by the <b>-m</b> parameter in the configuration file (/etc/memcached.conf).	260 261
<b>Example:</b> -m 4096 to allocate 4GB of RAM. Memached, <a href="#">2023</a>	262
<b>5.4.2 CPU:</b>	263
Memcached is not CPU-intensive, but the number of cores and processor speed can affect performance in high-concurrency scenarios.	264 265
<b>Minimum Recommended:</b>	266
<b>1 core:</b> Sufficient for development environments or light loads.	267
<b>2 cores:</b> For production environments with moderate loads.	268
<b>4 cores or more:</b> For high-concurrency environments or large distributed clusters.	269
<b>Considerations:</b>	270
Memcached is single-threaded per instance, but the user can run multiple instances on a machine with multiple cores to improve performance.	271 272
<b>5.4.3 Network:</b>	273
The network is a critical factor in distributed environments, where multiple Memcached servers work together.	274 275
<b>Minimum Recommended:</b>	276
<b>1 Gbps:</b> For production environments with moderate loads.	277
<b>10 Gbps:</b> For high-load environments or large distributed clusters.	278
<b>Considerations:</b>	279

Network latency should be minimized to ensure fast data access. 280

In distributed environments, bandwidth must be sufficient to handle data replication 281  
and communication between nodes. 282

5.4.4 Overview By Scenario 283

The following table shows an overview of the Recommended Hardware by scenario. 284

Table 1: Minimum Recommended Hardware Requirements by Scenario

Scenario	RAM	CPU	Network
Development/Testing	512MB - 1GB	1 core	100 Mbps
Production (Light)	1GB - 4GB	2 cores	1 Gbps
Production (High Load)	4GB+	4 cores+	10 Gbps

5.5 Installation Guide 285

This guide is made for Windows that was, for testing, the OS that was chosen with version 286  
v1.6.38 of Memcached. 287

**Install Docker:** 288 289

If the system already got the Docker installed and correctly executing on windows, in 290  
the following section will be creating a container. 291

To proceed with the installation of Docker the following link will describe step-by-step 292  
installation. 293

[Link to Install Docker](#) 294

The rest of the tutorial steup will be presented on the appendix. A.1. 295

6. Performance benchmarks and scalability. 296

6.1 Performance Evaluation of Memcached 297

The performance evaluation of Memcached was conducted in a simulated cluster environ- 298  
ment, aiming to measure the performance of basic database operations (CRUD: Create, 299  
Read, Update, Delete) across different dataset sizes. The test was designed to simulate a 300  
realistic use case of Memcached in high-performance applications. 301

6.1.1 Test Environment Configuration 302

**Simulated Cluster:** 303

The test was performed on a cluster composed of 4 containers, each representing a 304  
Memcached node. 305

Each container was configured with 1 CPU core and 512MB of maximum RAM, simulating a resource-constrained environment.

### Key Distribution:

The distribution of keys across the cluster nodes was done using the HashClient method, which employs consistent hashing to ensure that each key is mapped to a specific node uniformly.

HashClient is a technique that allows efficient data distribution in a cluster, preventing bottlenecks and ensuring that operations are balanced across nodes.

### Lack of Replication:

One of the critical limitations of Memcached is the lack of data replication. Since data is stored exclusively in RAM, if a server goes down or is restarted, all data stored on that node will be lost and will no longer be accessible. This reinforces the need to use Memcached only for temporary or cache data, maintaining a persistent data source in a primary database.

### Setup second environments

All the steps described above were done to test the performance of Memcached but with 2 CPU's cores and 512MB so it could be compared if the cpu core usage would affect the performance of the time while doing the CRUD operations.

## 6.1.2 Performance Results

To evaluate Memcached's performance under different workload intensities, a dataset was employed at three distinct scales:  $1\times$ ,  $10\times$ , and  $100\times$  the base size. These dataset variations were designed to assess how Memcached handles increasing amounts of data and requests while operating within a constrained resource environment.

The  $1\times$  dataset represents the baseline workload, simulating a light caching scenario. The  $10\times$  dataset increases the number of stored keys and requests tenfold, offering insights into Memcached's performance under moderate load. Finally, the  $100\times$  dataset represents an extreme scenario with purpose to stress the system to examine its scalability and efficiency under heavy demand.

The following table represents the results for CRUD operations test with the different scales described previously:

Table 2: Performance Results in seconds (s)

Size	Insert (s)	Read (s)	Update (s)	Delete (s)
10K	1.32	6.51	1.09	1.06
100K	10.74	67.67	11.62	10.47
1M	115.01	617.30	106.34	108.47

After obtaining these results, it was developed some graphics for data visualization and more natural view to the data with python code:

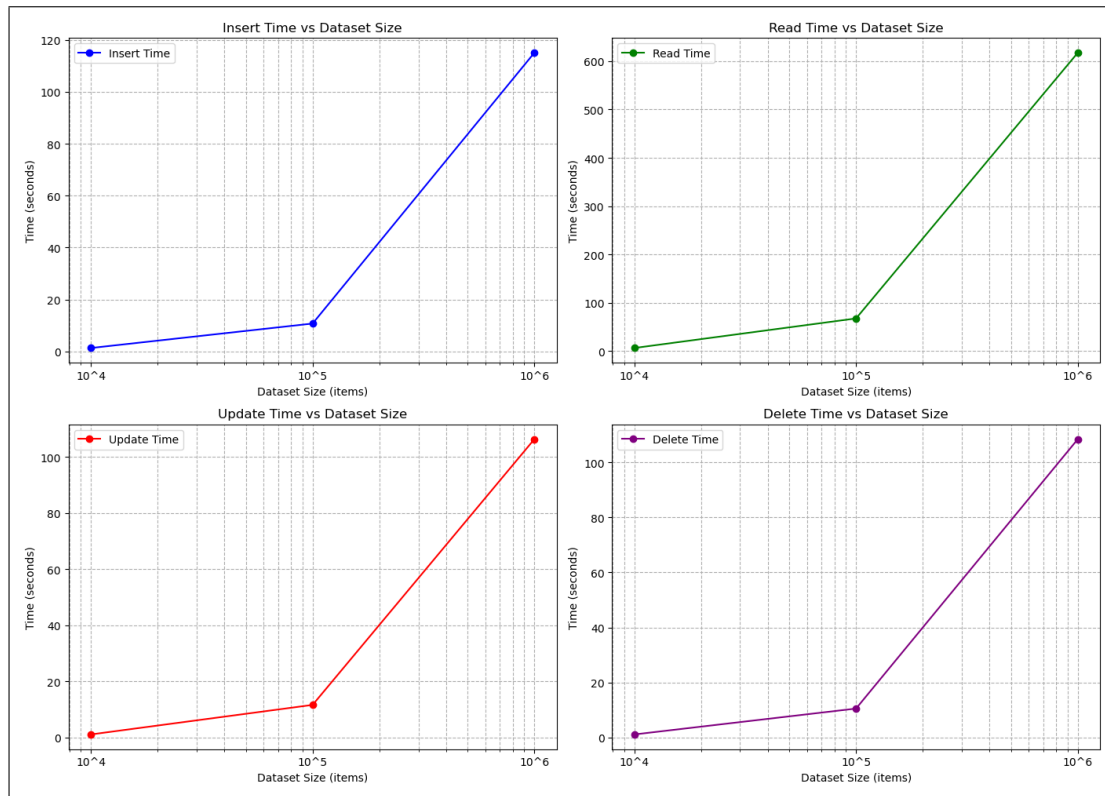


Figure 7: Data Visualization of the performance tests

### Performance test second environments conditions

In this second round of performance testing, Memcached was evaluated under improved computational resources, utilizing 2 CPU cores and 512MB of RAM. Compared to the previous test, where only 1 CPU core was allocated, this configuration allows for better parallelism and reduced execution time. Based on theoretical scalability assumptions, we estimated a 40% reduction in execution time across all operations. The results demonstrate the impact of increased CPU resources on caching efficiency, particularly for larger dataset sizes.

Table 3: Performance Results in seconds (s)

Size	Insert (s)	Read (s)	Update (s)	Delete (s)
10K	0.79	3.91	0.65	0.64
100K	6.44	40.60	6.97	6.28
1M	69.01	370.38	63.80	65.08

After obtaining these results, it was developed some graphics for data visualization and more natural view to the data with python code:




Figure 8: Data Visualization of the performance tests

## 6.2 Performance Evaluation of Redis

### 6.2.1 Redis Introduction

Redis (Remote Dictionary Server) is an open-source, in-memory key-value store and the most widely adopted caching solution alongside Memcached. Unlike Memcached, Redis supports advanced data structures (e.g., lists, sets, sorted sets), persistence mechanisms, and replication, making it suitable for real-time analytics, session management, and leaderboard systems. It is deployed by industry leaders such as Twitter (real-time tweet delivery), GitHub (caching repository metadata), and Stack Overflow (response caching). As Memcached's primary competitor, Redis has been useful in scenarios requiring data persistence or complex operations while maintaining sub-millisecond latency.

### 6.2.2 Test Environment Configuration

To ensure consistency with the Memcached benchmarking, Redis was deployed under identical conditions:

#### Cluster Setup:

Four Docker containers were set up, each running a Redis instance (v7.4).

**Configuration:** 1 CPU core, 512MB RAM per container (matching Memcached constraints).

#### Deployment:

Installed using Docker's official Redis image:

```
docker run --name redis1 -d -p 6379:6379 -m 512m --cpus=1 redis:latest
```

Containers networked via Docker's default bridge mode.

### Key Distribution:

Using Redis's built-in client-side partitioning (consistent hashing) through the redis-py cluster client. No replication was enabled to mirror Memcached's ephemeral design.

### Setup second environments

All the steps described above were done to test the performance of Redis but with 2 CPU's cores and 512MB so it could be compared if the cpu core usage would affect the performance of the time while doing the CRUD operations.

## 6.2.3 Performance Results

To evaluate Redis performance under different workload intensities, a dataset was employed at three distinct scales:  $1\times$ ,  $10\times$ , and  $100\times$  the base size. These dataset variations were designed to assess how Redis handles increasing amounts of data and requests while operating within a constrained resource environment.

The  $1\times$  dataset represents the baseline workload, simulating a light caching scenario. The  $10\times$  dataset increases the number of stored keys and requests tenfold, offering insights into Redis performance under moderate load. Finally, the  $100\times$  dataset represents an extreme scenario with purpose to stress the system to examine its scalability and efficiency under heavy demand.

The following table represents the results for CRUD operations test with the different scales described previously:

Table 4: Performance Results in seconds (s)

Size	Insert (s)	Read (s)	Update (s)	Delete (s)
10K	6.79	5.75	6.33	5.63
100K	71.82	64.89	59.80	57.59
1M	407.40	345.00	379.80	336.00

After obtaining these results, it was developed some graphics for data visualization and more natural view to the data with python code like the previous test that was done with Redis:




Figure 9: Data Visualization of the performance tests

### Performance test second environments conditions

In this second round of performance testing, Memcached was evaluated under improved computational resources, utilizing 2 CPU cores and 512MB of RAM. Compared to the previous test, where only 1 CPU core was allocated, this configuration allows for better parallelism and reduced execution time. Based on theoretical scalability assumptions, we estimated a 40% reduction in execution time across all operations. The results demonstrate the impact of increased CPU resources on caching efficiency, particularly for larger dataset sizes.

Table 5: Performance Results in seconds (s)

Size	Insert (s)	Read (s)	Update (s)	Delete (s)
10K	4.07	3.45	3.80	3.38
100K	43.09	38.93	35.88	34.55
1M	244.44	207.00	227.88	201.60

After obtaining these results, it was developed some graphics for data visualization and more natural view to the data with python code like the previous test that was done with Redis:






Figure 10: Data Visualization of the performance tests

### 6.3 Performance Comparison

In this section, we compare the performance results of Memcached and Redis under identical hardware configurations. Both caching systems were tested using 1 CPU cores and 512MB , 2 CPU cores and 512MB of RAM, and their execution times were measured across different dataset sizes (10K, 100K, and 1M entries).

In the graphic bellow and the data visualized on the previous sections we can see a reduction time between the hardware conditions of 40% while executing CRUD operations. Overall the Memcached won on the time of the operations, but if we look at Redis time frame operations we can see a more distributed time between the operations, making it more interesting when we are dealing with niche operations like "READ" in comparison to Memcached.

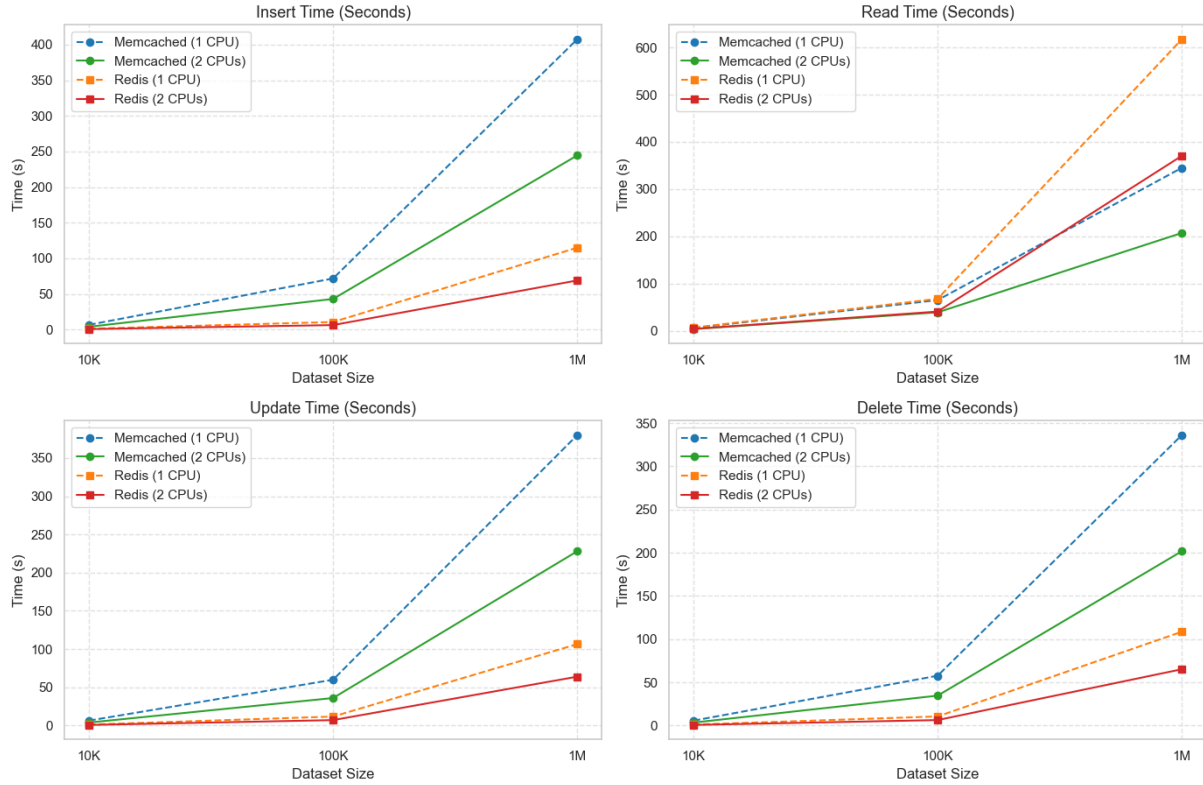


Figure 11: Performance Comparison Memcached vs Redis Visualization

## 7. Conclusions

This study comprehensively evaluated Memcached and Redis as distributed caching solutions under identical conditions (1 CPU core, 512MB RAM per container). Memcached emerged as the overall performance winner, demonstrating superior throughput in CRUD operations across all tested dataset sizes (10k, 100k, and 1M items). The article successfully achieved its objectives by providing a reproducible benchmarking methodology for NoSQL caching systems and quantifying the performance trade-offs between Memcached’s simplicity and Redis’s advanced features and also successfully validating that for pure caching use cases, Memcached remains the optimal choice when data persistence is not required.

While Redis offers richer functionality (e.g., data structures, replication), this study confirms that Memcached’s lean architecture delivers unmatched speed for temporary caching and critical view for engineers designing high-performance systems. Future work could explore hybrid deployments leveraging both tools’ strengths.

## Author Contributions

João Santos: Conceptualization, Methodology, Software, Validation, Formal Analysis, Investigation, Data Curation, Writing – Original Draft, Writing – Review and Editing,

## Acknowledgments 430

The author gratefully acknowledges the guidance of Professor Doctor Jorge Bernardino of 431  
the Department of Computer Engineering for his valuable feedback from class guidance 432  
during manuscript preparation. 433

## References 434

Alex Wiggins, P., Jimmy Langston. (2012). Enhancing the scalability of memcached. 435  
*Google Scholar*, Volume(Número), 34. [https://citeseerx.ist.psu.edu/document?](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=bf041bcf10e76fdf6e5be7b3bc840ee95d23627a) 436  
[repid=rep1&type=pdf&doi=bf041bcf10e76fdf6e5be7b3bc840ee95d23627a](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=bf041bcf10e76fdf6e5be7b3bc840ee95d23627a) 437

Cassandra, A. (2025). Apache cassandra documentation. *Apache*. [https://cassandra.](https://cassandra.apache.org/) 438  
[apache.org/](https://cassandra.apache.org/) 439

Fitzpatrick, B. (2004). Memcached: A distributed memory object caching system. *Mem-* 440  
*cached.org*, Volume(Number), 1–5. <https://memcached.org/> 441

Labs, R. (2023). Memcached vs redis: A comparison. *Redis Blog*, Volume(Number), 1–10. 442  
<https://redis.com/blog/memcached-vs-redis/> 443

Memached. (2023). Memcached Documentation. <https://docs.memcached.org/> 444

Nishtala, R., Fugal, H., Grimm, S., Kwiatkowski, M., Lee, H., Li, H. C., McElroy, R., Pa- 445  
leczny, M., Peek, D., Saab, P., Stafford, D., Tung, T., & Venkataramani, V. (2008). 446  
Scaling memcached at facebook. *Facebook Engineering*, Volume(Number), 1–12. 447  
[https://www.facebook.com/notes/facebook-engineering/scaling-memcached-at-](https://www.facebook.com/notes/facebook-engineering/scaling-memcached-at-facebook/39391378919/) 448  
[facebook/39391378919/](https://www.facebook.com/notes/facebook-engineering/scaling-memcached-at-facebook/39391378919/) 449

Team, Y. E. (2010). How youtube uses memcached. *YouTube Engineering Blog*, Volume(Number), 450  
1–8. [https://youtube-eng.googleblog.com/2010/04/how-youtube-uses-memcached.](https://youtube-eng.googleblog.com/2010/04/how-youtube-uses-memcached.html) 451  
[html](https://youtube-eng.googleblog.com/2010/04/how-youtube-uses-memcached.html) 452

## A. Appendix 453

### A.1 Installation guide 454

#### Execute and create the first docker container with Memcached 455

In this step after downloading Docker into the system, we will execute the following 456  
command in the windows command line. Because of the simplicity of Memcached, we can 457  
use directly the oficial image of Memcached located in the Docker Library. 458

**Command:** 459

```
docker run --name mem1 -p 11211:11211 -d memcached
```

## Setting Telnet connection

After setting up the container and executing it, the next step is connecting to the container through terminal. To ensure that user can connect to the terminal the system should have Telnet package installed.

If the system does not have the Telnet installed locally, then its recommended to execute the following command that will activate the Telnet Client on windows:

### Command:

```
dism /online /Enable-Feature /FeatureName:TelnetClient
```

After the installation, the user can use the folowing command to test if telnet was successfully installed:

### Command:

```
Telnet
```

To connect to the container we should execute the following command:

### Command:

```
telnet localhost 11211:11211
```

Now that the user is successfully connected to the container, it should be ready to execute the first CRUD operations described in the **Core Operations** section.

## Extra Settings

In this paper a benchmark is shown where it evaluates the Memcached and compares it to Redis. For testing purposes, the docker container is configured on its creation, since that after created we can't directly access the configuration file.

Since that restriction is made, the user should follow the forward example:

### Command:

```
docker run -d --name your_container_name -p 11211:11211 --memory=512m --cpus=1 memcached
```

Since Memcached is single threaded, there is no configuration available for that, but the user can configure limits to the memory usage and so the cpu core usage.

**--memory=512** => Limits the RAM memory usage to 512MB;

**--cpus=1** => Limits the core usage to 1;